

Consera Deep learning course 4 Note

CNN

Week 1

Vertical edge detection

$$\begin{array}{c}
 \begin{array}{|c|c|c|c|c|c|} \hline
 3 & 0 & 1 & - & - & - \\ \hline
 - & 5 & 8 & - & - & - \\ \hline
 2 & 7 & 1 & - & - & - \\ \hline
 - & - & - & - & - & - \\ \hline
 \end{array} &
 \begin{array}{c}
 \text{"convolution"} \\
 \downarrow \\
 \ast
 \end{array} &
 \begin{array}{|c|c|c|} \hline
 1 & 0 & 7 \\ \hline
 1 & 0 & -1 \\ \hline
 1 & 0 & -1 \\ \hline
 \end{array} &
 = &
 \begin{array}{|c|c|c|c|} \hline
 -5 & 0 & 8 & - \\ \hline
 0 & 2 & 2 & 3 \\ \hline
 -1 & -1 & -1 & -1 \\ \hline
 \end{array} \\
 6 \times 6 & & 3 \times 3 & & 4 \times 4
 \end{array}$$

$$3 \times 1 + 1 \times 1 + 2 \times 1 + 0 \times 0 \\
 + 5 \times 0 + 7 \times 0 + 1 \times -1 + 8 \times -1 + 1 \times -1 = -5$$

↙ vertical edges

$$\begin{array}{c}
 \begin{array}{ccccccccc}
 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\
 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\
 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\
 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\
 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0
 \end{array} &
 \begin{array}{c}
 \ast \\
 1 & 0 & -1 \\
 1 & 0 & -1 \\
 1 & 0 & -1
 \end{array} &
 = &
 \begin{array}{ccccccccc}
 0 & 3 & 0 & 3 & 0 & 0 & 0 & 0 & 0 \\
 0 & 3 & 0 & 3 & 0 & 0 & 0 & 0 & 0 \\
 0 & 3 & 0 & 3 & 0 & 0 & 0 & 0 & 0 \\
 0 & 3 & 0 & 3 & 0 & 0 & 0 & 0 & 0
 \end{array}
 \end{array}$$

6×6

3×3

4×4

$$\begin{array}{c}
 \begin{array}{ccc}
 1 & 0 & -1 \\
 2 & 0 & -2 \\
 1 & 0 & -1
 \end{array} &
 \begin{array}{c}
 3 & 0 & -3 \\
 1 & 0 & -10 \\
 3 & 0 & -3
 \end{array}
 \end{array}$$

sobel filter

Scharf filter

$w_1 w_2 w_3$
 $w_4 w_5 w_6$
 $w_7 w_8 w_9$

} can be
trained

Padding

$$\begin{array}{c} 6 \times 6 \\ n \times n \\ \xrightarrow{\quad \quad} \\ \begin{array}{l} 3 \times 3 \\ f \times f \\ n-f+1 \times n-f+1 \\ 6-3+1=4 \end{array} \end{array} \longrightarrow 4 \times 4$$

$$\begin{array}{c} \text{Padding} \\ \swarrow \\ \begin{array}{c} 6 \times 6 \rightarrow 8 \times 8 \\ n \times n \quad \begin{array}{l} n+2p \\ \times \\ n+2p \end{array} \end{array} \end{array}$$
$$\begin{array}{c} 3 \times 3 \\ f \times f \\ n+2p-f+1 \times n+2p-f+1 \\ 8-3+1=6 \end{array} \longrightarrow 6 \times 6$$

"Valid" convolution $n \times n \times f \times f \rightarrow n-f+1 \times n-f+1$

usually odd

"Same" convolution $(n+2p-f+1) \times (n+2p-f+1)$

Strided convolution

$$7 \times 7 \times 3 \times 3 = 3 \times 3$$

$$\begin{array}{c} \text{stride} = 2 \\ n \times n \times f \times f \\ p=0 \quad s=2 \end{array} \qquad \qquad \qquad \begin{array}{c} \downarrow \text{floor} \\ = \left[\frac{n+2p-f}{s} + 1 \right] \times \left[\frac{n+2p-f}{s} + 1 \right] \\ 3 \times 3 \end{array}$$

Convolution on RGB images

$$6 \times 6 \times 3 \quad * \quad 3 \times 3 \times 3 = 4 \times 4$$

height width channels height width channels height width

P.G.B

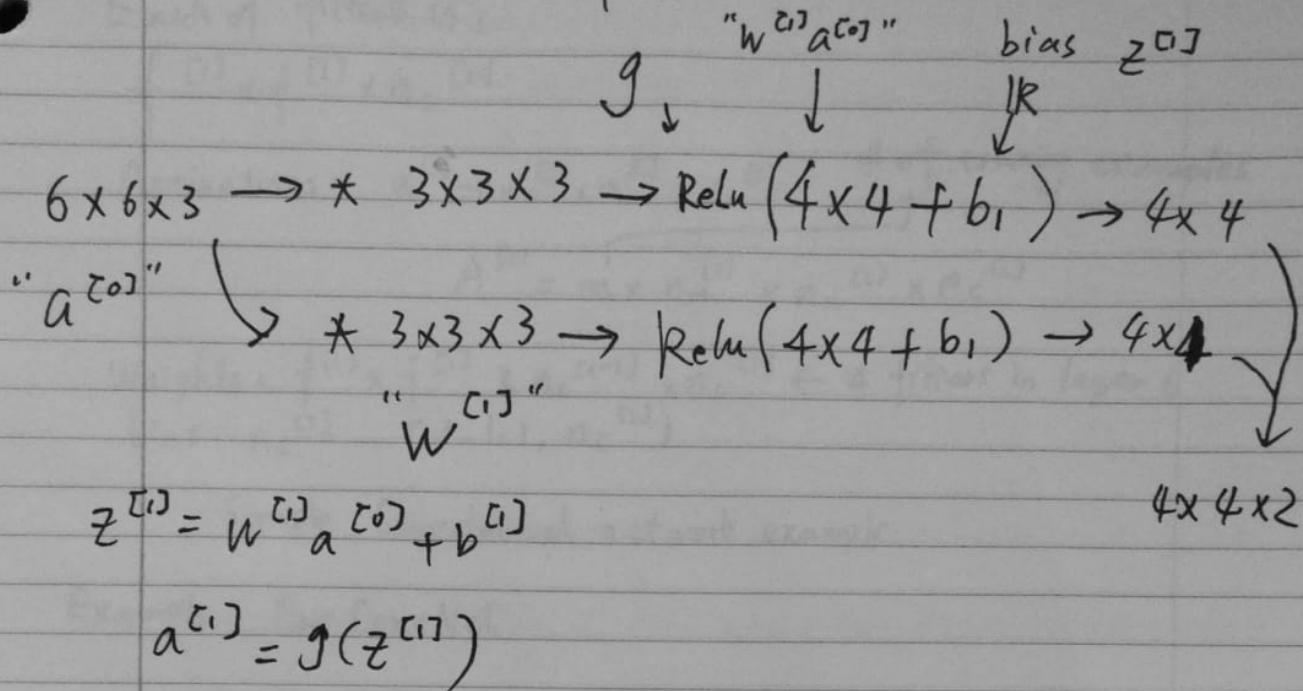
$$27 \text{ numbers} * 27 \text{ numbers} = 1 \text{ number}$$
$$(9+9+9) \quad R \quad G \quad B$$

100
1000
10000

multiple filter

$$\begin{array}{ccc}
 \text{multiple filter} & \xrightarrow{\quad\quad\quad} & \text{vertical edge} \\
 \nearrow * & & 3 \times 3 \times 3 = 4 \times 4 \\
 6 \times 6 \times 3 & & \text{horizontal edge} \\
 \nearrow * & & \xrightarrow{\quad\quad\quad} 4 \times 4 \times 2 \\
 & & 3 \times 3 \times 3 = 4 \times 4
 \end{array}$$

One layer of convolution network



e.g.: 10 filters that are $3 \times 3 \times 3$ in one layer
How many parameters does that layer have?

$$\begin{aligned} & 3 \times 3 \times 3 \\ & 27 \text{ parameters} \\ & + \text{bias} \\ & = 28 \text{ parameters} \times 10 = 280 \text{ parameters} \end{aligned}$$

Summary of notation

$$\begin{aligned} f^{[l]} &= \text{filter size} & \text{Input: } n_H^{[l-1]} \times n_W^{[l-1]} \times n_C^{[l-1]} \\ l &= \text{layer number} & \text{Output: } n_H^{[l]} \times n_W^{[l]} \times n_C^{[l]} \\ s^{[l]} &= \text{stride} & n^{[l]} = \left\lfloor \frac{n^{[l-1]} + 2P^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rfloor \\ n^{[l]} &= \text{num of filters} \end{aligned}$$

Each of filters is:

$$f^{[l]} \times f^{[l]} \times n_c^{[l]}$$

Activations: $a^{[l]} \rightarrow n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$ # of training examples
 $A^{[l]} = m \times \overbrace{n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}}$

Weights: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]} \leftarrow$ # filters in layer l
bias: $n_c^{[l]} - (1, 1, 1, n_c^{[l]})$

Simple Convolutional network example

Example Box ConvNet

$$\begin{array}{ll} 39 \times 39 \times 3 & \xrightarrow{f^{[0]}=3} 37 \times 37 \times 10 \xrightarrow{f^{[1]}=5} 17 \times 17 \times 20 \\ n_H^{[0]} = n_W^{[0]} = 39 & s^{[0]}=1 \quad n_H^{[1]} = n_W^{[1]} = 37 \quad n_H^{[2]} = n_W^{[2]} = 17 \\ n_c^{[0]} = 3 & p^{[0]}=0 \quad n_c^{[1]} = 10 \quad s^{[1]}=2 \quad n_c^{[2]} = 20 \\ n_c^{[0]} = 10 & p^{[1]}=0 \quad n_c^{[2]} = 20 \end{array}$$

$$\begin{array}{l} \xrightarrow{f^{[2]}=5} 7 \times 7 \times 40 \xrightarrow{\text{softmax}} \hat{y} \\ s^{[2]}=2 \quad n_H^{[2]} = n_W^{[2]} = 7 \\ n_c^{[2]} = 40 \quad 1960 \rightarrow 1 \end{array}$$

$n_H, n_W \downarrow$

$n_c \uparrow$

Types of layer

Convolution (CONV)

Pooling (POOL)

Fully Connected (FC)

Pooling layers

Max pooling: take max of numbers in a region

$$4 \times 4 \rightarrow 2 \times 2$$

$f=2$
 $s=2$ → parameters do not change in gradient decent

$$5 \times 5 \times 2 \xrightarrow{f=3} 3 \times 3 \times 2$$

$s=1$

Average pooling: take average of numbers in a region

note: max pooling is used more often than average pooling

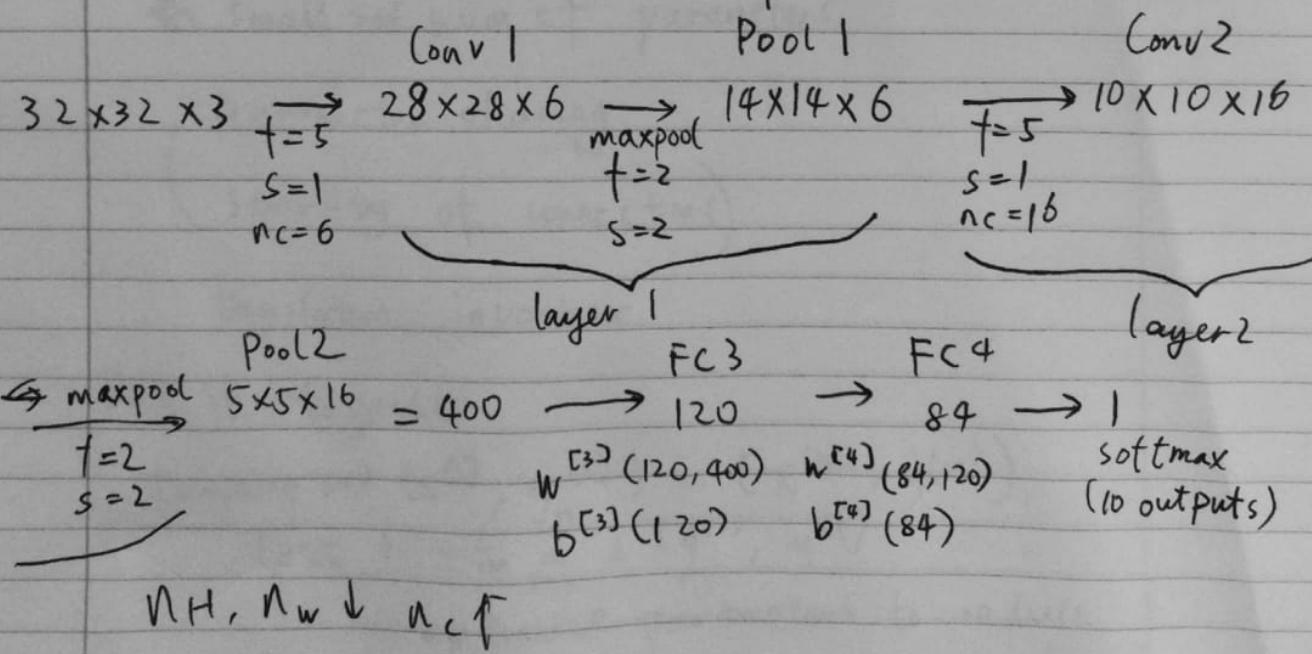
Hyperparameters

f , s , max or average pooling

No parameters to learn

$$\begin{aligned} n_H &\times n_w \times n_c \\ &\downarrow \\ &[\frac{n_H-f}{s}+1] \times [\frac{n_w-f}{s}+1] \\ &\times n_c \end{aligned}$$

Neural network example (Le Net - 5)



Shape

	Activation shape	Activation Size	# parameters
Input	(32, 32, 3)	$32 \times 32 \times 3$ ↓ 3072 $a^{(0)}$	0
CONV1 ($t=5, s=1$)	$\frac{32+0-5+1}{2} = 28$ (28, 28, 8)	$28 \times 28 \times 8$ ↓ 6272	$5 \times 5 \times 8 + 8$ 208
POOL1	$\frac{28}{2} = 14$ (14, 14, 8)	$14 \times 14 \times 8$ ↓ 1568	no parameters to learn 0
CONV2 ($t=5, s=1$)	$\frac{14+0-5+1}{2} = 10$ (10, 10, 16)	$10 \times 10 \times 16$ ↓ 1600	$5 \times 5 \times 16 + 16$ 416
POOL2	$\frac{10}{2} = 5$ (5, 5, 16)	$5 \times 5 \times 16$ ↓ 400	no parameters to learn 0
FC 3	(120, 1)	120	$400 \times 120 + 1$ 4800
FC 4	(84, 1)	84	$120 \times 84 + 1$ 1008
Softmax	(10, 1)	10	$84 \times 10 + 1$ 84

Why convolutions

* Small set num of parameters

(parameter sharing
sparsity of connections)

translation invariance

Put it together

Training set $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$

$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)})$$

... use parameters to reduce |

Classic networks:

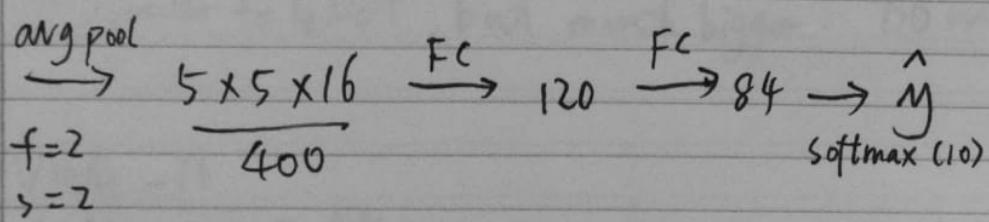
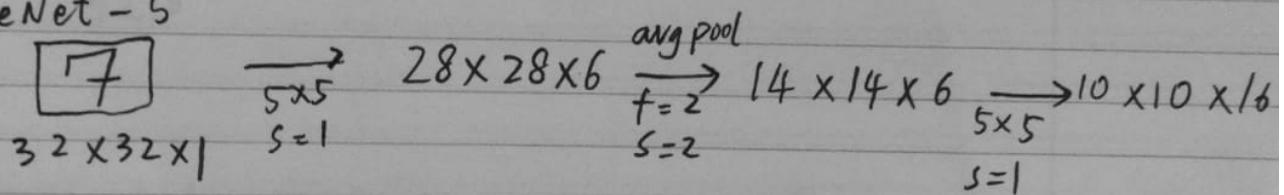
- LeNet - 5
- AlexNet
- VGG

ResNet (152)

Inception

Case studies classic networks

LeNet - 5



60k parameters

$n_H, n_w \downarrow n_C \uparrow$

conv pool conv pool fc fc output

AlexNet

$$227 \times 227 \times 3 \xrightarrow[11 \times 11]{S=4} 55 \times 55 \times 96 \xrightarrow[\substack{3 \times 3 \\ S=2}]{\text{max-pool}} 27 \times 27 \times 96 \xrightarrow[5 \times 5]{\text{Same}} 27 \times 27 \times 256$$

$$\xrightarrow[\substack{3 \times 3 \\ S=2}]{\text{max-pool}} 13 \times 13 \times 256 \xrightarrow[\text{Same}]{3 \times 3} 13 \times 13 \times 384 \xrightarrow[\text{same}]{3 \times 3} 13 \times 13 \times 384 \xrightarrow[\text{same}]{3 \times 3} 13 \times 13 \times 256$$

$$\xrightarrow[\substack{3 \times 3 \\ S=2}]{\text{max-pool}} \frac{6 \times 6 \times 256}{9216} = 9216 \xrightarrow{\text{FC}} 4096 \xrightarrow{\text{FC}} 4096 \xrightarrow{\text{softmax}} 1000$$

- Similar to LeNet, but much bigger: 60 m parameters

VGG - 16

CONV = 3×3 filters, $S=1$, same MAX-POOL = 2×2 , $S=2$

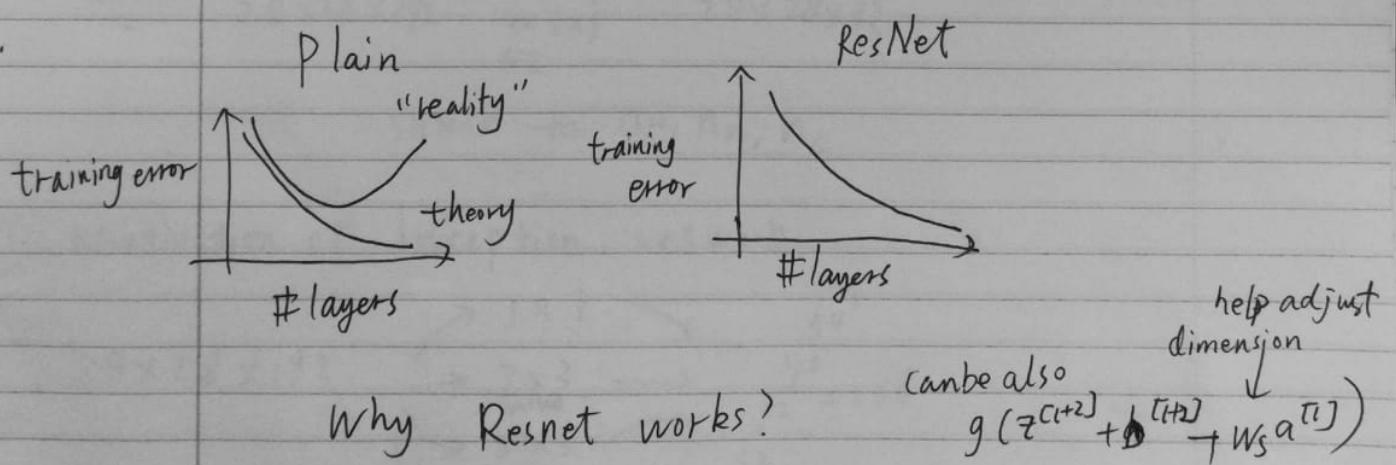
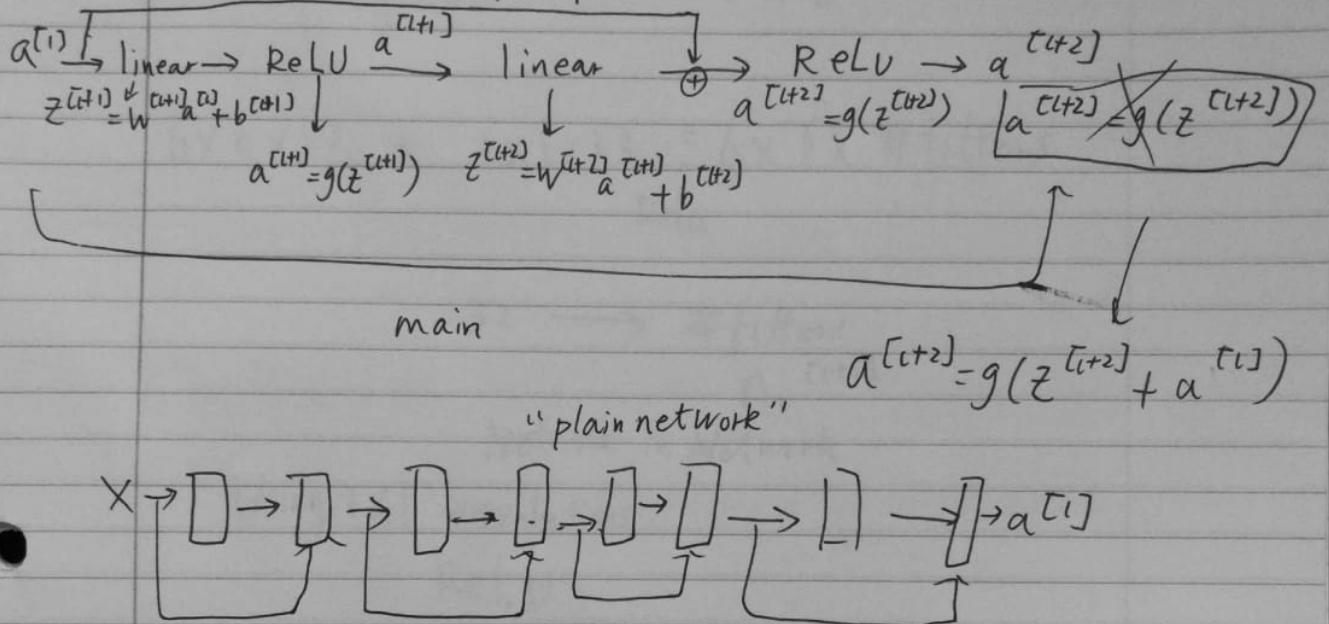
$$224 \times 224 \times 3 \xrightarrow[\substack{\text{CONV} 64 \\ \times 2}]{\text{CONV} 64} 224 \times 224 \times 64 \xrightarrow{\text{POOL}} 112 \times 112 \times 64 \xrightarrow[\substack{\text{CONV} 128 \\ \times 2}]{\text{CONV} 128} 112 \times 112 \times 128$$

$$\xrightarrow{\text{POOL}} 56 \times 56 \times 128 \xrightarrow[\substack{\text{CONV} 256 \\ \times 3}]{\text{CONV} 256} 56 \times 56 \times 256 \xrightarrow{\text{POOL}} 28 \times 28 \times 256 \xrightarrow[\substack{\text{CONV} 512 \\ \times 3}]{\text{CONV} 512} 28 \times 28 \times 512$$

$$\xrightarrow{\text{POOL}} 14 \times 14 \times 512 \xrightarrow[\substack{\text{CONV} 512 \\ \times 3}]{\text{CONV} 512} 14 \times 14 \times 512 \xrightarrow{\text{POOL}} 7 \times 7 \times 512 \xrightarrow[4096]{\text{FC}} 4096 \xrightarrow{\text{FC}} 1000 \xrightarrow{\text{softmax}}$$

(38 m parameters, 16 weight layers)

Residual block (help train very deep network)
 "short cut"/skip connection



Why Resnet works?

$$a^{[l+2]} = g(z^{[l+2]} + b^{[l+2]} + w_s a^{[l]})$$

where $z^{[l+2]} = h^{[l+2]} a^{[l+1]}$

and $b^{[l+2]} = \alpha > 0$

make sure doesn't hurt performance

1×1 convolution

$$6 \times 6 * \text{⊗} [1 \times 1] = 6 \times 6$$

$$6 \times 6 \times 32 * 1 \times 1 \times 32 = 6 \times 6 \times \# \text{filters}$$

ReLU

$$32 \longrightarrow \# \text{filters}$$
$$n_c^{[l+1]}$$

Network in Network

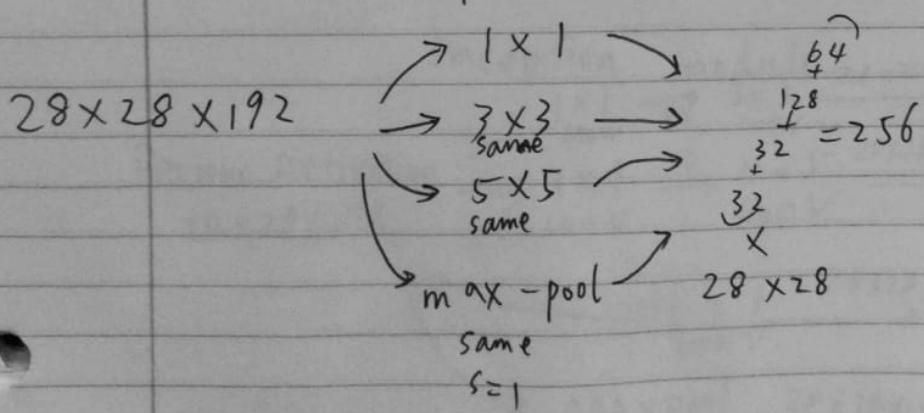
Using 1×1 convolutions

ReLU

$$28 \times 28 \times 192 \xrightarrow[\text{32}]{\text{conv } 1 \times 1} 28 \times 28 \times 32$$

shrink to n_h, n_w, n_c

Motivation of Inception network



The problem of computational cost

$$28 \times 28 \times 192 \xrightarrow{\text{conv}} 28 \times 28 \times 32$$

5x5,
 same
 32

32 filter

Computation:

$$\underbrace{28 \times 28 \times 32}_{\text{32 filter}} \times \underbrace{5 \times 5 \times 192}_{\text{32 filter}} = 120M$$

Using 1×1 convolution, "bottleneck layer"

$$28 \times 28 \times 192 \xrightarrow[\text{1x1, } 1^b]{\text{conv}} 28 \times 28 \times 16 \xrightarrow[\text{5x5, } 5^b]{\text{conv}} 28 \times 28 \times 32$$

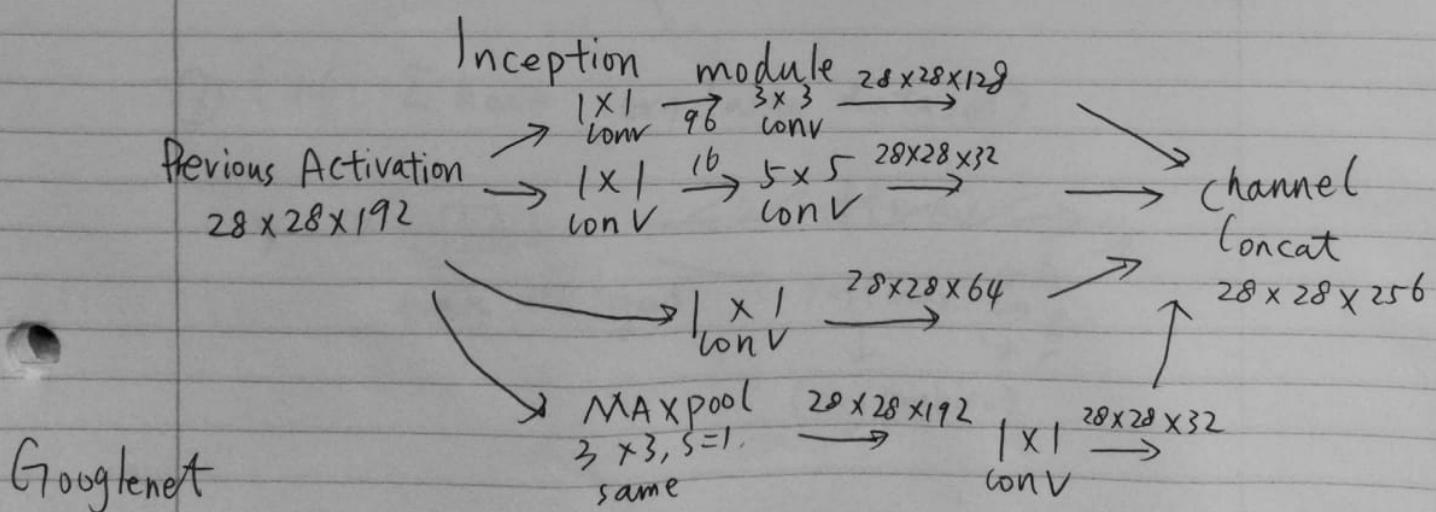
\downarrow
 $5 \times 5 \times 16$

Computation:

$$28 \times 28 \times 16 \times 192 = 2.4M$$

$$+ 28 \times 28 \times 32 \times 5 \times 5 \times 16 = 10.0M$$

$$= 12.4M$$

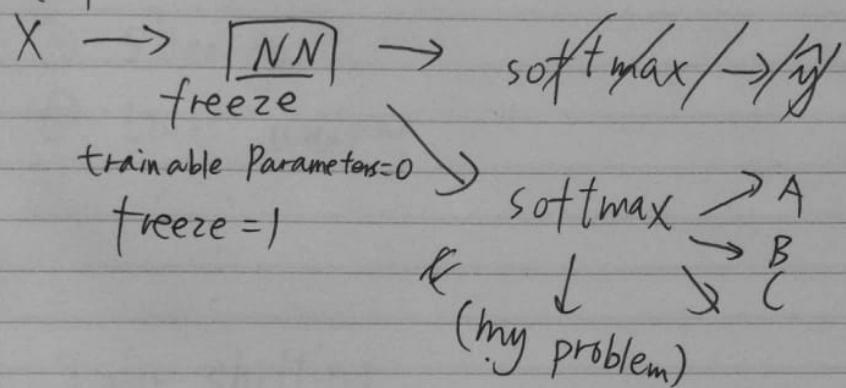


advice for using convNets

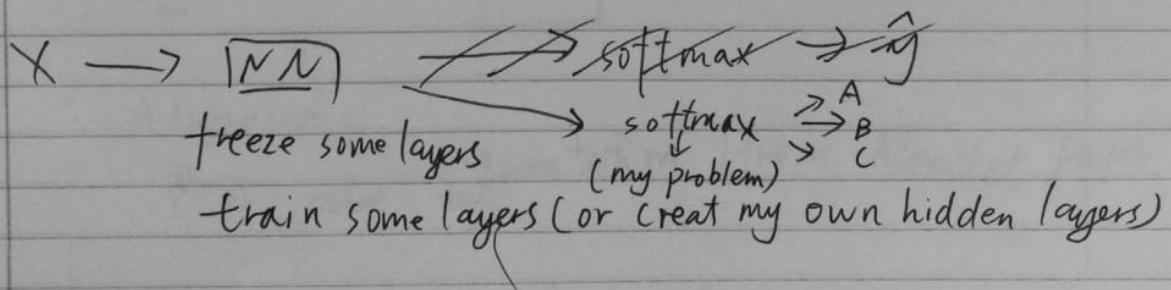
1. Using open-source implementation

2. transfer learning

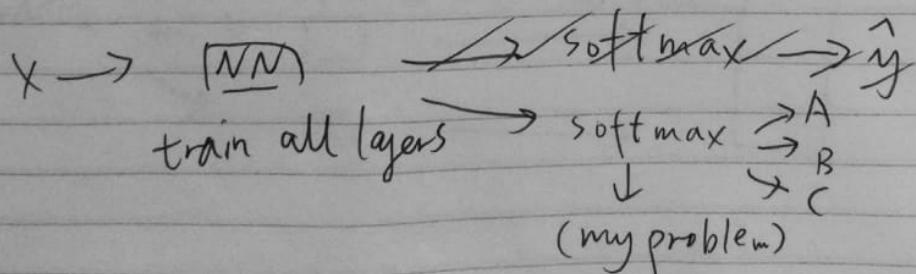
① (if I have small dataset)



② (if I have many pictures)



③ (if I have very large dataset)



3. Data augmentation

① Mirroring

② Random Cropping

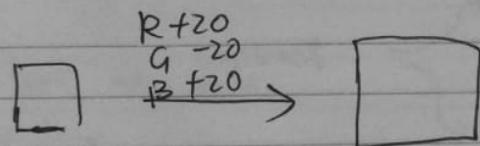
③ Rotation

④ Shearing

⑤ Local warping

:

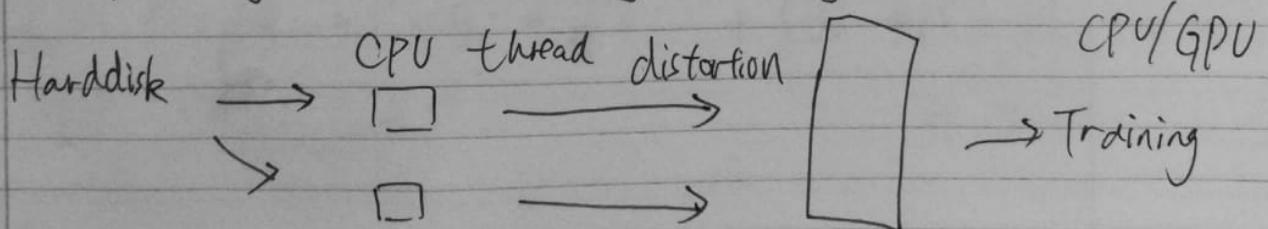
Color shifting



Advanced:

PCA color augmentation (from AlexNet paper)

Implementing distortions during training



Tips for doing well on benchmarks
Ensembling (not used in production)
(Average different models' output)

Multip

Multi-crop at test time

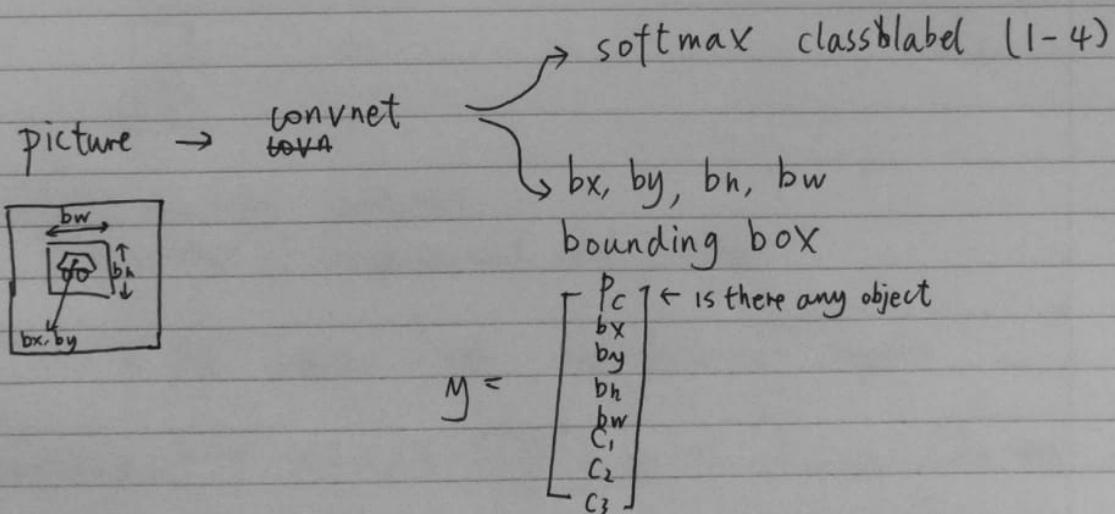
Classification with localization

1. pedestrian

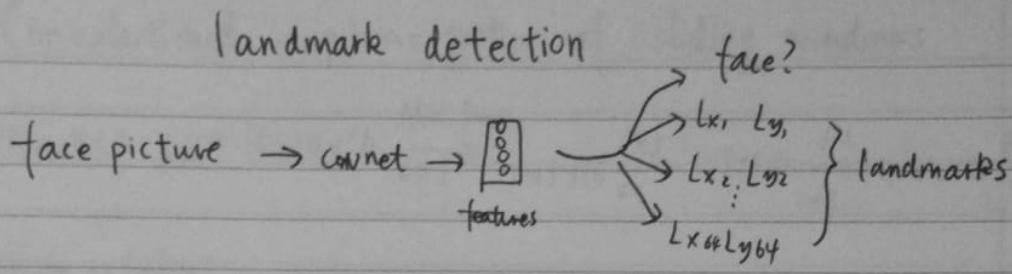
2. car

3. motorcycle

4. background



$$L(\hat{y}, y) = \begin{cases} (\hat{y}_1 - y_1)^2 + (\hat{y}_2 - y_2)^2 + \dots + (\hat{y}_8 - y_8)^2 & \text{if } y_1 = 1 \\ (\hat{y}_1 - y_1)^2 & \text{if } y_1 = 0 \end{cases}$$



e.g.



• black points are landmarks

Object detection

Car detection example

Training set y



1

\rightarrow convnet



0

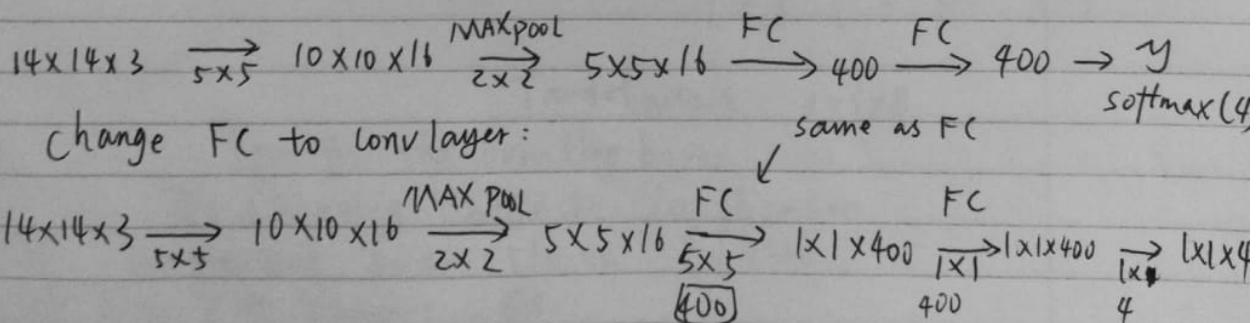


1

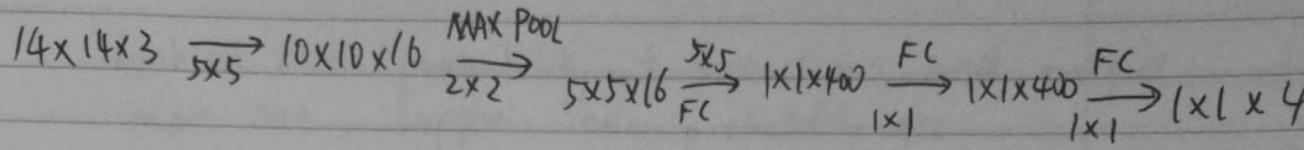
Sliding windows detection

Disadvantage: computational cost, slow

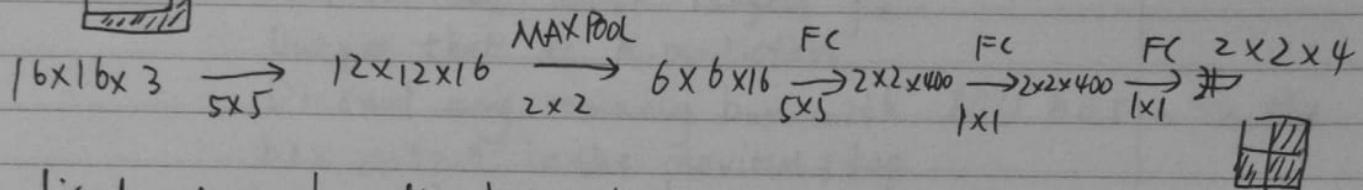
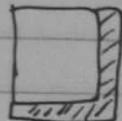
Turning FC layer into convolutional layers



Convolutional implementation of sliding windows

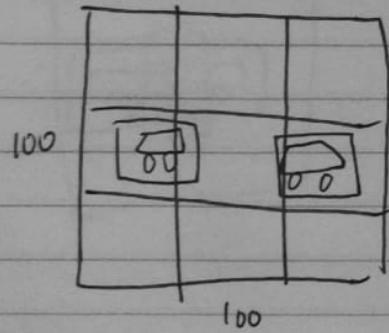


add stripe as follows



disadvantage: bounding box not too accurate

Output accurate bounding box
YOLO algorithm



Labels for training
For each grid cell

$$y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix} \begin{bmatrix} [] \\ [] \\ [] \\ [] \\ [] \\ [] \\ [] \\ [] \end{bmatrix}$$

Target output $3 \times 3 \times 8$

Specify the bounding boxes

Evaluating object localization

Intersection over Union (IOU)



= size of intersection

size of

"correct" if IOU ≥ 0.5

Non-max suppression example

look at the probability of detections, and output max one, suppress lower probability ones

Discard all boxes with $p_c \leq 0.6$

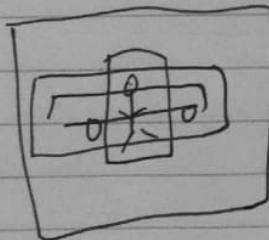
while there are many remaining boxes

- Pick the box with largest p_c

Output that as a prediction

- Discard any remaining box with $IOU \geq 0.5$ with the box output in the previous step

Anchor & boxes



$$\text{define } y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ \vdots \\ c_1 \\ c_2 \\ c_3 \end{bmatrix} \left\{ \begin{array}{l} \text{anchor box 1} \\ \vdots \\ \text{anchor box 2} \end{array} \right.$$

object is assigned to grid cell that contains object's mid point and anchor box which has highest IOU

Region proposal

R-CNN

Propose regions

of sliding window

Fast R-CNN Use convolution implementation to propose regions ✓ propose

Faster R-CNN Use convolution network to propose regions

Face verification vs. face recognition

Verification

- Input image, name/ID
- Output whether the input image is that of the claimed person

Recognition

- Has a database of k persons
- Get an input image
- Output ID if the image is any of the k persons (or "not recognized")

One-shot learning

Learn from one example to recognize the person again

$d(\text{img1}, \text{img2})$ = degree of difference between images

if $\begin{cases} d \leq \tau & \text{"same"} \\ d > \tau & \text{"different"} \end{cases}$ } verification

Siamese network

$X^{(1)} \rightarrow \text{CNN} \rightarrow \cancel{\text{softmax}}$
 $+ f(X^{(1)})$ "encoding of $X^{(1)}$ "

$X^{(2)} \rightarrow \text{CNN} \rightarrow f(X^{(2)})$ "encoding of $X^{(2)}$ "

$$d(X^{(1)}, X^{(2)}) = \|f(X^{(1)}) - f(X^{(2)})\|_2^2$$

Learning Objective

(A)	(P)	(A)	(N)
Anchor	Positive	Anchor	negative

d to be large small d to be large

$$\underbrace{\|f(A) - f(P)\|^2}_{d(A,P)} \leq \underbrace{\|f(A) - f(N)\|^2}_{d(A,N)}$$

$$\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + \alpha \leq 0$$

Loss function

margin \rightarrow make sure
 $f(\text{img}) = \vec{0}$
 condition ~~is~~ covered

Given 3 images A, P, N

$$L(A, P, N) = \max(\underbrace{\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + \alpha}_{> 0}, 0)$$

$$J = \sum_{i=1}^m L(A^{(i)}, P^{(i)}, N^{(i)})$$

Training set: 10k pictures of 1k persons
 (Note: there may be more than 1 pic for one person)

Choosing the triplets of A, P, N
 if A, P, N are chosen randomly,
 $d(A, P) + \alpha \leq d(A, N)$ is easily satisfied

Choose triplets that're "hard" to train on
 $d(A, P) + \alpha \leq d(A, N)$

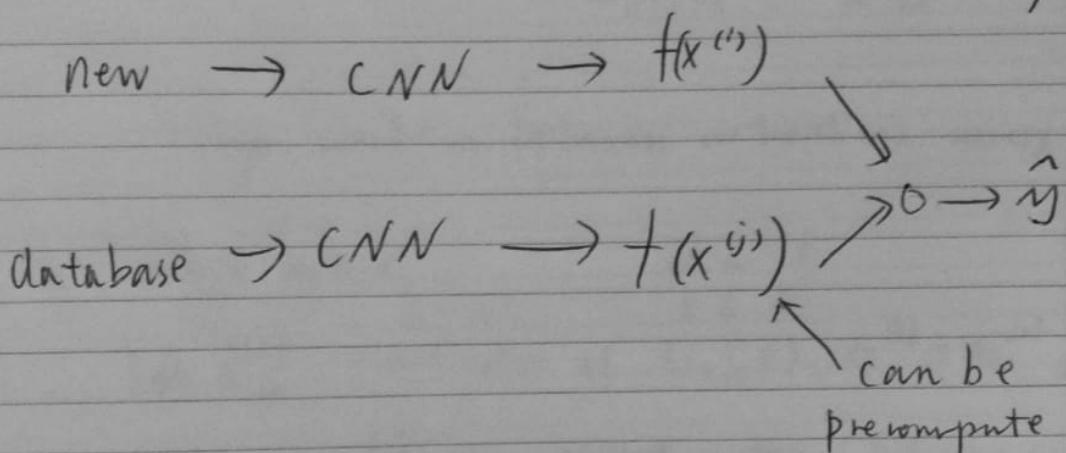
$$d(A, P) \approx d(A, N)$$

Verification (as binary classification)

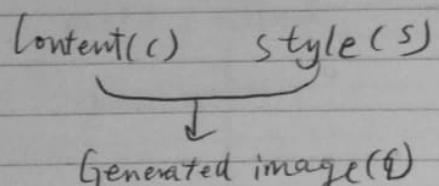
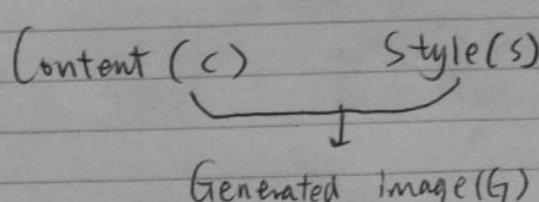
Learning the similarity function

$$\hat{y} = \sigma \left(\sum_{k=1}^{128} w_i \underbrace{|f(x^{(i)})_k - f(x^{(i)})_k|}_{\frac{(f(x^{(i)})_k - f(x^{(i)})_k)^2}{f(x^{(i)})_k + f(x^{(i)})_k}} + b \right)$$

or $\left(\frac{(f(x^{(i)})_k - f(x^{(i)})_k)^2}{f(x^{(i)})_k + f(x^{(i)})_k} \right) \xrightarrow{\text{can be precompute}} \hat{x}^2 \right)$



Neural style transfer



lost function

$$J(G) = \alpha J_{\text{content}}(C, G) + \beta J_{\text{style}}(S, G)$$

Find the generated image G

1. Initiate G randomly

G: 100 x 100 x 3

2. Use gradient descent to minimize J(G)

$$G = G - \frac{\alpha}{\partial G} J(G)$$

Content cost function

$$J(G) = \alpha J_{\text{content}}(C, G) + \beta J_{\text{style}}(S, G)$$

~~choose~~: layer L : $a^{[L](c)}$, $a^{[L](G)}$ are the activations

$$J_{\text{content}}(C, G) = \frac{1}{2} \|a^{[L](c)} - a^{[L](G)}\|^2$$

style: correlation between activations across channels

style matrix

H w c
↓ ↓ ↓

Let $a_{ijk}^{[l]}$ = activation at (i, j, k) . $G^{[l]}$ is $n_c^{[l]} \times n_c^{[l]}$

$G_{kk'}^{[l]} \leftarrow$ compare correlations between k and k'

$$G_{kk'}^{[l]} = \sum_i^H \sum_j^W a_{ijk}^{[l]} a_{ijk'}^{[l]} \quad \begin{matrix} \leftarrow \text{style image} \\ \uparrow \quad \uparrow \end{matrix} \quad \begin{matrix} \text{also called} \\ \text{"gram matrix"} \end{matrix}$$

$$G_{kk'}^{[l](G)} = \sum_i^{n_H^{[l]}} \sum_j^{n_W^{[l]}} a_{ijk}^{[l](G)} a_{ijk'}^{[l](G)}$$

$$\begin{aligned} J_{\text{style}}^{[l]}(S, G) &= \frac{1}{(2n_H^{[l]} n_W^{[l]})^{n_c^{[l]}}} \|G_{kk'}^{[l](S)} - G_{kk'}^{[l](G)}\|_F^2 \\ &= \frac{1}{(2n_H^{[l]} n_W^{[l]})^{n_c^{[l]}}} \sum_{k'} \left(G_{kk'}^{[l](S)} - G_{kk'}^{[l](G)} \right)^2 \end{aligned}$$

$$J_{\text{style}}(S, G) = \sum_l \lambda^{[l]} J_{\text{style}}^{[l]}(S, G)$$

Convolutions in 2D and 1D

$$2D \\ [4 \times 14 \times 3 \times 5 \times 5 \times 3] \rightarrow [10 \times 10 \times 16]$$

1D

$$[4 \times 1] * [5 \times 1] \rightarrow [10 \times 16]$$

3D

$$N \times W \times D \times C \\ [4 \times 14 \times 14 \times 1] * [5 \times 5 \times 5 \times 1] \rightarrow [10 \times 10 \times 10 \times 16] \\ 16 \text{ filter} \\ N_{C+1}$$